

## Simple Example

Let's now look at a simple example that includes a number of rules and a variety of facts in a working memory. We'll use a subset of the ZOOKEEPER rules first proposed by Patrick Henry Winston, the Ford Professor of the Artificial Intelligence and Computer Science Department at the Massachusetts Institute of Technology [\[Winston 1993\]](#).

Consider the following rules shown in [Listing 8.1](#).

### Listing 8.1: Sample Problem Adapted from Winston's ZOOKEEPER Example.


```
(defrule bird-test
  (has-feathers ?)
=>
  (bird ?)
)

(defrule mammal-test
  (gives-milk ?)
=>
  (mammal ?)
)

(defrule ungulate-test1
  (mammal ?)
  (chews-cud ?)
=>
  (is-ungulate ?)
)

(defrule ungulate-test2
  (mammal ?)
```

```
(has-hoofs ?)
=>
  (is-ungulate ?)
)
```



Now consider that our working memory contains the following facts:

```
(gives-milk animal)
```

```
(has-hoofs animal)
```

We begin with the match phase where we try to match the facts in our working memory with the antecedents in our rule set. Neither fact matches the antecedents of the first rule (`bird-test`), but an antecedent is matched in the second test (`mammal-test`). We save this rule off into our conflict set and continue our search through the remaining rules. No other rules match given the current working memory, so in this case we have the simple case of a conflict set containing one rule. Since there is no conflict, the `mammal-test` rule is permitted to fire (the action phase), resulting in the working memory as shown below:

```
(gives-milk animal)
```

```
(has-hoofs animal)
```

```
(mammal animal)
```

We begin again at the match phase and walk through the rules looking for matching antecedents. In this case, our conflict set contains two rules, `mammal-test` and `ungulate-test2`. Conflict resolution is also simple given this case, because only one rule here has any affect on the working memory (`ungulate-test2`). The first rule has previously fired; therefore, nothing new can be added to our

working memory and the rule can be omitted from the conflict set. Moving on to the action phase and firing the ungulate-test2 rule results in the following working memory:

(gives-milk animal)

(has-hoofs animal)

(mammal animal)

(ungulate animal)

From this very simple example, our rules system has deduced that the animal in question is an ungulate (a hoofed animal). It reasoned this after first determining that the animal was a mammal (knowing that it gave milk) and then using this information, along with the hoof information, determined that the animal was an ungulate.

Conflict resolution in the example shown included two of the most basic mechanisms for determining which rule to fire. Other mechanisms could be provided for scenarios that are more complicated. For example, the conflict resolver could pick the rule to fire from the conflict set that had the largest number of antecedents. This provides for the most complicated case and could help the system reason to the most relevant goal state and ignore the simpler cases, which might lead the system down unnecessary paths.


## Sample Application

We'll now look at another example that provides a form of fault tolerance within an embedded domain. While reviewing this application, we'll also present the other consequent primitives that are available.

### Fault Tolerance

In a very simple application, we'll encode rudimentary knowledge of sensor management into our rules. A redundant set of sensors exists from which one can be set as the active sensor (and thus be used by the larger system). Either of the sensors may be working (or failed), but only one may be active at any given time. If neither of the sensors is working, then neither can be active. Additionally, our mode must change to denote that a problem exists within the sensor subsystem.

In this application, it is assumed that we're operating in a blackboard architecture (see [Figure 8.3](#)). Blackboard architecture contains a number of agents that use and produce data on the blackboard. The blackboard serves as a working space where facts are communicated amongst the agents. Agents are triggered by data on the blackboard and can manipulate the data (adding, changing, or removing), which may trigger other agents to continue the work process.



**Figure 8.3:** Graphical depiction of a blackboard architecture.

The blackboard architecture is interesting because it not only provides the ability for agents to communicate with one another, but also to coordinate and synchronize their activities.

## Rules Definition

The rules file for our sample fault tolerance demonstration is shown in [Listing 8.2](#). This set of rules introduces a number of new consequents (actions) that can be performed.

### Listing 8.2: Fault Tolerance Demonstration Rules File.

```
(defrule init
  (true null)                                ; antecedent
=>
  (add (sensor-active none))                  ; consequents
  (add (sensor-working sensor1))
  (add (sensor-working sensor2))
  (add (mode failure))
  (enable (timer 1 10))
  (print ("default rule fired!"))
  (disable (self))
)

;
; Define active rule-set
;

(defrule sensor-failed
  (sensor-working ?)
  (sensor-failed ?)
=>
  (delete (sensor-working ?))
)

(defrule check-active
  (sensor-active ?)
  (sensor-failed ?)
=>
  (delete (sensor-active ?))
  (add (sensor-active none))
)
```

```
)
```

```
(defrule make-working
  (sensor-active none)
  (sensor-working ?)
=>
  (add (sensor-active ?))
  (delete (mode failure))
  (add (mode normal))
  (delete (sensor-active none))
)
```

```
(defrule failure
  (mode normal)
  (sensor-active none)
  (sensor-failed sensor1)
  (sensor-failed sensor2)
=>
  (add (mode failure))
  (delete (mode safe))
  (delete (mode normal))
)
```

```
; Use triggers to simulate timed events...
```

```
(defrule trigger1
  (timer-triggered 1)
=>
  (print ("Sensor 1 failure.\n"))
  (add (sensor-failed sensor1))
  (enable (timer 2 10))
  (delete (timer-triggered 1))
)
```

```
(defrule trigger2
  (timer-triggered 2)
=>
  (print ("Sensor 2 failure.\n"))
  (add (sensor-failed sensor2))
  (enable (timer 3 10))
  (delete (timer-triggered 2))
)
```

```

)

(defrule trigger3
  (timer-triggered 3)
=>
  (print ("Sensor 1 is now working.\n"))
  (delete (sensor-failed sensor1))
  (add (sensor-working sensor1))
  (enable (timer 4 10))
  (delete (timer-triggered 3))
)

(defrule trigger4
  (timer-triggered 4)
=>
  (print ("Sensor 2 is now working.\n"))
  (delete (sensor-failed sensor2))
  (add (sensor-working sensor2))
  (enable (timer 1 10))
  (delete (timer-triggered 4))
)

```

This example contains nine rules, but only four are actual operating rules. The first rule (`init`) in [Listing 8.2](#) is an initialization rule. Note that the antecedent is (`true null`), which always resolves to true. This initialization rule permits us to seed the working memory with an initial set of facts (as defined by the `add` commands). The `enable` command allows us to enable a timer, which when fired can trigger another rule. We'll use timers in this example to perform simulated events from the environment. These would be events (facts) placed into the working memory from other agents operating within the environment (recall [Figure 8.3](#)). [Note: The timers used in this example are modeled after Microsoft's *The Age of Kings*, which uses a rules file to define the game AI.] The `timer` command uses two arguments, the numeric id of the timer and the number of seconds until the timer should fire (in this case, timer 1 will fire in 10 seconds). The `print` command is a debugging command that provides us with some visibility into the reasoning of the system. The final command in this rule, `disable`, is very important. It provides the ability to disable the rule from the set of available rules. This rule can then never be fired again, an obvious need for rules whose antecedents are always matched (such as the `init` rule).

Once the init rule has fired, the working memory will contain:

```
(sensor-active none)
(sensor-working sensor1)
(sensor-working sensor2)
(mode failure)
```

Rule **make-working** would be the next rule to fire (using sensor1 as the matched parameter), leaving working memory as:

```
(sensor-working sensor1)
(sensor-working sensor2)
(sensor-active sensor1)
(mode normal)
```

At this point, no further rules can fire and the system remains in this state until the timer fires (after 10 seconds). When a timer fires, the system simply adds the fact to working memory **((timer-triggered 1))**. This new fact in working memory then causes the rule to fire that handles the particular timer (using the trigger event as the rule's antecedent). The event triggered by the timer simulates the failure of sensor1, modifying working memory as follows:

```
(sensor-working sensor1)
(sensor-working sensor2)
(sensor-active sensor1)
(mode normal)
(sensor-failed sensor1)
```

Rules **sensor-failed** and **check-active** then fire, leaving working memory in the following state:


```
(sensor-working sensor2)
(mode normal)
(sensor-failed sensor1)
(sensor-active none)
```

Finally, rule **make-working** fires, leaving working memory as:

```
(sensor-working sensor2)
(mode normal)
(sensor-failed sensor1)
(sensor-active sensor2)
```



At this point, the system is operational with a working sensor from the redundant pair. The rules also provide for no working sensors with an indication of this event using the mode fact. The derivation of the remaining rule firings is left for the reader.

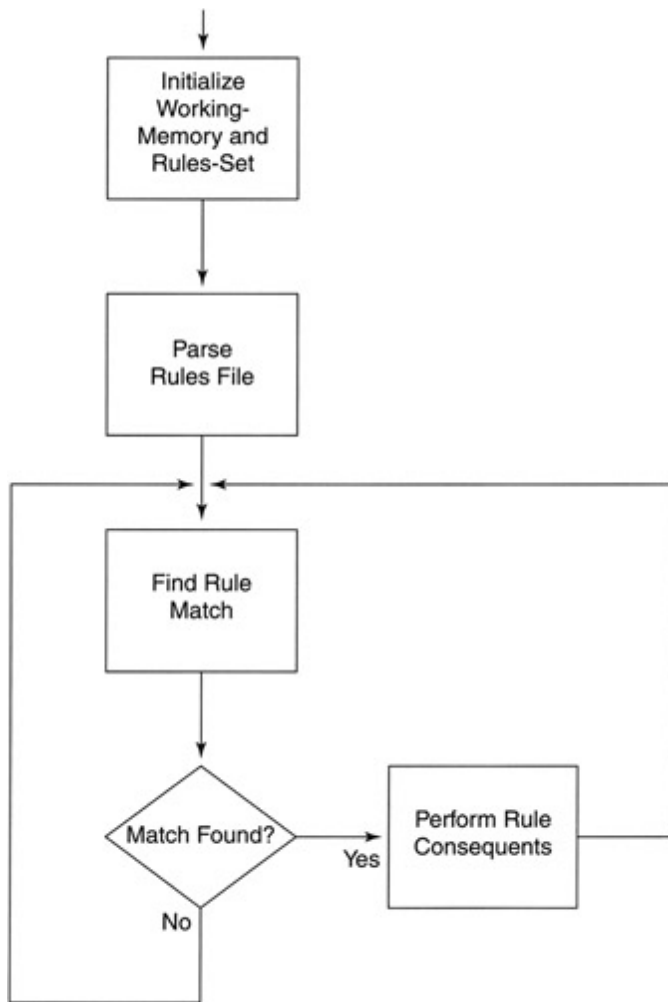
**Note** This example can be performed with the rules-based system application on the CD-ROM. The fault-tolerance example shown above is contained in the file  *fault.rbs*.

**On the CD** The source code for the rules based system can be found on the CD-ROM at ./software/ch8/rbs/.

## Source Discussion

Let's now look at a simple implementation of a rules-based system using the rules and fact format discussed thus far.

The basic flow of the software is shown in [Figure 8.4](#). After initializing (clearing) the working memory and the rules-set, the rules file specified by the user is parsed and the rules loaded into the rules-set. The software then goes into a loop trying to match a rule given the current working memory. When a rule is matched, the execution logic is called to perform the consequents of the particular rule. Whether or not a rule was fired, the software continues by trying to match another rule. This is the basic simplified flow of the rules-based system.



**Figure 8.4:** Basic flow of the rules-based system.

Let's first look at the data structures that are used by the source ([Listing 8.3](#)).

### **Listing 8.3: Rules-Based System Data Structures.**

```
#define MEMORY_ELEMENT_SIZE    80
```

```
#define MAX_MEMORY_ELEMENTS    40
```